Final Exam Answers – CS 343 Fall 2015

Instructors: Peter Buhr and Ashif Harji

December 18, 2015

These are not the only answers that are acceptable, but these answers come from the notes or lectures.

1. (a) **2 marks** *Staleness* occurs when a task reads data that is temporally in the past with respect to a current data value.

    Freshness occurs when a task reads data that is temporally in the future with respect to a current data value.

   (b) **3 marks**

   - Not if signalled task must implicitly reacquire the mutex lock before continuing.
   - ⇒ signaller must release the mutex lock.
   - There is now a race between signalled and calling tasks, resulting in *barging*.

   (c) **1 mark** No value. The baton is fictitious.

   (d) **1 mark** Multiple threads must acquire resources in the same order to prevent deadlock.

   (e) **1 mark** A cycle in a process graph, does not imply a deadlock if any resource has multiple instances.

   (f) **2 marks** The preemption victim must be restarted (but where) or killed.

2. (a) **6 marks**

   ```
   1   shared list/queue<int> q; // shared queue of integers

       producer task                        consumer task

   1   region q {                       region q {
   1     await q.size() < 10;       1     await ! q.empty();
   1       // add item to queue      1       // remove item from queue
       }                                }
   ```

   (b) **2 marks** For signal the signalling task continues execution until it waits or exits, and the signalled task is delayed (on the Acceptor/Signaller stack).

    For signalBlock the signalling task is delayed (on the A/S stack), and the signalled task continues execution until it waits or exits.

   (c) **4 marks** If a writer is using the resource (wcnt > 0), achieve cooperation by preventing all calls except endWrite, after which the writer has finished using the resource.

    If a reader is using the resource (rcnt > 0), achieve cooperation by preventing all calls except N calls to endRead, after which all readers have finished using the resource.

   (d) **2 marks** There is a fixed (bounded) number of restarted tasks that check their predicates before new tasks can enter.

   (e) **2 marks** Barging occurs when a calling task acquires the monitor *before* a signalled task.

   (f) **1 mark** Spurious wakeup means a task blocked on a lock can unblock without being released or signalled.

3. (a) **2 marks** The long form of the _Accept statement is necessary to add _When clauses and action statements onto individual _Accept clauses OR

   _When(C1) Accept(m1) S1; **or** _When(C2) Accept(m2) S2;

   (b) **2 marks** It is important to have as little code as possible in the mutex members of a server task to increase concurrency of the client because the client spends as little time as possible performing server work.

   The mutex-member code goes into the task main of the server, so the server performs the administration work rather than the client.

   (c) **2 marks** The client has to be unblocked so it can raise the exception on its stack (rather than server's stack).

   (d) **2 marks** A task accepts its destructor to know when to terminate.

   Accepting the destructor is unusual because it behaves as a signal rather than a signalBlock, i.e., the acceptor runs first and then the acceptee (destructor) runs.

   (e) **1 mark** The general approach to increase client-side is turning synchronous calls into asynchronous calls.

   (f) **2 marks** The advantage of futures is that they do not have an explicit protocol because the server pretends to instantaneously return a computed a result.

   (g) **2 marks** The _Select statement allows a task to block waiting for multiple futures to become available depending on a complex relationships among the futures.

   _Select(f1 || f2 && f3);

4. (a) **3 marks** False sharing is when threads are accessing disjoint (non-shared) variables but the variables are actually shared on the same cache line.

   The sharing causes *cache bouncing* if each thread is writing to the variables.

   Aligning or separating (padding) the variables so they are on separate cache lines.

   (b) **2 marks** Lazy cache-consistency occurs because caches cannot be update instantly, allowing readers to see own write before acknowledgement so other threads continue to read stale data.

   (c) **1 mark** Declaration qualifier **volatile** prevents variables from being hidden in registers.

   (d) **1 mark** No, if a programming language does not have a memory model, concurrency cannot be implemented with a library approach.

   (e) **2 marks** advantage: lock-free has no deadlock OR hold resource on preemption

   disadvantage: not general to arbitrary critical section OR performance questionable.

   (f) **2 marks** Ada requeue mechanism is not as powerful as internal scheduling, because internal scheduling preserves data and execution state while blocked but requeue does not.

   (g) **2 marks** Go uses channels to support direct communication. Go uses a select statement to choose among a number of channels for data or block until data arrives.

5. (a) **2 marks** An address space is an area of memory addressed at 0, so addresses cannot be used among address spaces, which it the precursor to separate memories in a distributed system.

   (b) **1 mark** There is a problem passing *integer* values between different computers because the byte ordering of the integers OR big versus little endian.

   (c) **2 marks** One thread is designated the *root* and its sends and receives occur as written but all other threads do the opposite action (receive or send).

   (d) **2 marks** Message passing is used to carry the call arguments through the network to the server and carry the results back through the network to the client.

6. (a) i. **3 marks**

```
        V()
        1   counter += 1;
        P()
        1   if ( counter == 0 ) _Accept( V );
        1   counter -= 1;
```

ii. **3 marks**

```
        1   uCondition bench;
        V()
        -   counter += 1;
        1   bench.signal();
        P()
        1   if ( counter == 0 ) bench.wait();
        -   counter -= 1;
```

iii. **6 marks**

```
        1   unsigned int tickets = 0, serving = counter; // prevent bargers
        V()
        1   serving += 1;
        1   signalAll();
        P()
        1   unsigned int myticket = tickets;            // select ticket
        1   tickets += 1;                               // advance ticket for next barger
        1   while ( myticket >= serving ) wait();       // my turn to proceed ?
```

iv. **3 marks**

```
        1   AUTOMATIC_SIGNAL;
        V()
        -   counter += 1;
        1   RETURN();
        P()
        1   WAITUNTIL( counter > 0, , );
        -   counter -= 1;
```

(b) **2 marks** Yes, it works. The P-waiter leaves first, and then the V-signaller, but that order is unimportant.

3

7. **23 marks**

```
    void fillShuttle( unsigned int noOfClients ) {
1       shuttles.signalBlock();                                    // get shuttle id
1       for ( unsigned int i = 0; i < noOfClients; i += 1 ) {
1           clients.front().delivery( shuttleId );
1           clients.pop_front();
        } // for
1       numClientsWaiting -= noOfClients;
    } // Coordinator::fillShuttle

    void main() {
1       for ( ;; ) {
1           _Accept( ~Coordinator ) {
1               break;
1           } or _Accept( timeUp ) {
1               if ( ! shuttles.empty() && numClientsWaiting > 0 ) fillShuttle( numClientsWaiting );
1           } or _Accept( checkIn ) {
1               if ( numClientsWaiting >= ShuttleSize ) fillShuttle( ShuttleSize );
1           } or _Accept( getRide ) {
1               if ( ! shuttles.empty() && numClientsWaiting >= ShuttleSize ) fillShuttle( ShuttleSize );
            } // _Accept
        } // for

1       shuttingDown = true;
        // can shutdown in any order
1       for ( unsigned int i = 0; i < NumClients; i += 1 ) {
1           if ( clients.empty() ) _Accept( getRide );
1           clients.front().exception( new Closed );
1           clients.pop_front();
        } // for
1       for ( unsigned int i = 0; i < NumShuttles; i += 1 ) {
1           if ( shuttles.empty() ) _Accept( checkIn );
1           shuttles.signalBlock();
        } // for
1       _Accept( timeUp );
    } // Coordinator::main
```