

# Floating Point – Error and Stability

CS370 - May 6, 2016

# Who are you?

1. What area(s) of CS are you most interested in?
2. What aspect/topic of the course are you most interested in?
3. What are your career plans/hopes for after graduation?

# Floating Point – Quick Recap

We approximated the real numbers with finite storage by limiting precision (# of digits,  $t$ ) and range (exponent,  $p$ ):

$$0.d_1d_2 \dots d_t \times \beta^p$$

An FP system  $F$  guarantees that:

- $fl(x) = x(1 + \delta)$  for  $x \in \mathbb{R}$  with  $|\delta| \leq E$ .
- $a \oplus b = fl(a + b) = (a + b)(1 + \delta)$  for  $a, b \in F$  and  $|\delta| \leq E$ .

These rules let us bound error in our calculations.

# Error Bounds and Condition Number

Last time, we said that  $(a \oplus b) \oplus c$  satisfies:

$$E_{rel} \leq \frac{|a| + |b| + |c|}{|a + b + c|} (2E + E^2)$$

The term  $\frac{|a|+|b|+|c|}{|a+b+c|} \dots$

- describes how the error  $E$  is magnified along the way.
- may be called a *condition number* for this calculation.

# Cancellation Errors

When is  $\frac{|a|+|b|+|c|}{|a+b+c|}$  large? When is it small?

Worst case magnification when denominator is small.

i.e., when  $|a + b + c| \ll |a| + |b| + |c|$ .

This occurs when quantities have differing signs and similar magnitudes, leading to cancellation.

# Error Bound Example #1

$$\begin{aligned}a &= 2000 \\b &= -3.234 \\c &= -2000 \\F &= \{10, 4, -10, 10\}\end{aligned}$$

Consider  $(a \oplus b) \oplus c$ :

True result:  $-3.234$     FP result:  $-3.000$

Error bound:

$$E_{rel} \leq \frac{|a|+|b|+|c|}{|a+b+c|} (2E + E^2) \approx \frac{4003.234}{3.234} \cdot 2 \left( \frac{1}{2} 10^{-3} \right) \approx 1.238.$$

$E^2$  term is very small, so we drop it.

A rather weak bound; large potential relative error.

## Error Bound Example #2

$$\begin{aligned}a &= 2000 \\b &= -3.234 \\c &= -2000 \\F &= \{10, 4, -10, 10\}\end{aligned}$$

Consider  $(a \oplus c) \oplus b$ :

True result:  $-3.234$     FP result:  $-3.234$

Observation #1: Re-ordering operations often changes the results.

Observation #2: Our error bound is the *same as before*.

The condition number only gives a *worst-case* bound.

$$F = \{10, 4, -10, 10\}$$

## Error Bound Example #3

Consider  $(a \oplus b) \oplus c$  for  $a = -2000, b = -3.234, c = -2000$ :  
(i.e. all matching signs now).

True:  $-4003.234$     Computed:  $-4003$

Condition number:  $\frac{|a|+|b|+|c|}{|a+b+c|} = \frac{4003.234}{4003.234} = 1.$

Error bound:  $E_{rel} \leq 1 \cdot (2E + E^2) \approx 2E = 10^{-3}.$

Observation: Avoiding cancellation gives better bounds on error.



# Catastrophic Cancellation Error

*Catastrophic* cancellation occurs when subtracting numbers of the same magnitude, *and* the input numbers contain error.

e.g.,

$$\begin{array}{ccccccc} 173.00063 & - & 173.00017 & = & 0.00046 \\ \underbrace{\hspace{1.5cm}} & & \underbrace{\hspace{1.5cm}} & & \underbrace{\hspace{1.5cm}} \\ \text{Correct} & & \text{Correct} & & \text{Erroneous} \\ \text{digits} & & \text{digits} & & \text{digits} \end{array}$$

All *significant* digits cancelled out; the result might have no correct digits whatsoever!

# Benign Cancellation

However, if input quantities are known to be *exact*, we have our usual guarantee on floating point ops (addition, subtraction, etc.):

$$w \oplus z = fl(w + z) = (w + z)(1 + \delta).$$

# Round-off Errors We've Seen

Adding large and small numbers (very different magnitudes).

- Smaller digits get lost or “swamped”!
- Rule of thumb: Try to sum numbers of approximately same size.

Subtracting nearby numbers *that contain error*.

- Loss of accuracy due to catastrophic cancellation.
- Rule of thumb: Try to reformulate computations to avoid cancellation.

# Taylor series example, revisited

So what's the reason we observed that

$$e^x \approx 1 + x + \frac{x^2}{2!} + \frac{x^3}{3!} + \dots$$

performs so much worse than



$$e^{-x} = \frac{1}{e^x} = \frac{1}{1 + x + \frac{x^2}{2!} + \frac{x^3}{3!} + \dots}$$

as an algorithm to evaluate  $e^{-5.5}$  in FP?

Latter *avoids alternating signs* in the sums which lead to cancellation.

# A Visualized Example of Cancellation Error

Compare:  $f(x) = \frac{1 - \cos^2 x}{x^2}$  and  $g(x) = \frac{\sin^2 x}{x^2}$  near  $x = 0$ . True solution approaches 1 as  $x \rightarrow 0$ .




MATLAB plot with IEEE double precision (i.e.,  $\approx 15 - 17$  decimal digits).

# FP Error: A Geometric Application Example

Do two line segments intersect?

A useful ***robust*** test for this must consider round-off error!

Both false positives or false negatives can occur.



e.g. “Epsilon geometry”  
[Guibas et al. 1989]

# Sources of Error

When solving a problem numerically, there are many possible error sources:

- Round-off error – due to FP representation and arithmetic.
- Truncation error – eg. truncating a Taylor series after  $n$  terms.
- Uncertainty/error in the input itself (e.g. from measurements).
- Error/approximation in our mathematical model of the “real” problem.

We will (continue to) focus largely on the first two.

# Conditioning of Problems

Problems may be *ill-conditioned* or *well-conditioned*.

For problem  $P$ , with input  $I$  and output  $O$ , if a change to the input,  $\Delta I$ , gives a “small” change in the output  $\Delta O$ ,  $P$  is *well-conditioned*.

Otherwise,  $P$  is *ill-conditioned*.

This is a property of the problem, *independent* of any specific implementation (algorithm or hardware).

Conditioning is relative; there’s a “sliding scale”.



# Stability of an *Algorithm*

If any initial error in the data is magnified by an *algorithm*, the algorithm is considered numerically *unstable*.

Can lead to meaningless results, for *seemingly* reasonable methods on reasonable problems.

# Conditioning v.s. Stability

Conditioning of a problem:

- How sensitive is the *problem* itself to errors/changes in input?

Stability of a numerical algorithm:

- How sensitive is the *algorithm* to errors/changes in input?

Note that:

1. An *algorithm* can be unstable even for a well-conditioned problem!
2. An ill-conditioned problem limits how well we can expect an algorithm to perform.

# Stability Analysis of an *Algorithm*

Consider the integration problem

$$I_n = \int_0^1 \frac{x^n}{x + \alpha} dx$$

for a given  $n$  where  $\alpha$  is some fixed parameter.

The course notes gives a recursive algorithm, for  $n \geq 0$ :

$$I_0 = \log \frac{1 + \alpha}{\alpha}, \quad I_n = \frac{1}{n} - \alpha I_{n-1}.$$

# Stability Analysis of an *Algorithm*

Hence that recurrence algorithm is *unstable* for solving  $I_n = \int_0^1 \frac{x^n}{x+\alpha} dx$  for values of  $|\alpha| > 1$ .

e.g. if you code up this recurrence in C / Matlab / etc.:

$$I_{100} = 6.64 \times 10^{-3} \text{ for } \alpha = 0.5 \quad \text{Correct!}$$

$$I_{100} = 2.1 \times 10^{22} \text{ for } \alpha = 2.0 \quad \text{Wrong!}$$

# Summary of FP

- $F$  is not  $\mathbb{R}$ !
- We can use **round-off error analysis** to bound the errors incurred by floating point operations.
- We can analyze whether initial errors grow or shrink to determine the **stability** of algorithms.

# Further reading on FP [Optional]

“What Every Computer Scientist Should Know About Floating-Point Arithmetic”, by David Goldberg, 1991.

## Appendix D

### What Every Computer Scientist Should Know About Floating-Point Arithmetic

**Note** – This appendix is an edited reprint of the paper *What Every Computer Scientist Should Know About Floating-Point Arithmetic*, by David Goldberg, published in the March, 1991 issue of Computing Surveys. Copyright 1991, Association for Computing Machinery, Inc., reprinted by permission.

#### Abstract

Floating-point arithmetic is considered an esoteric subject by many people. This is rather surprising because floating-point is ubiquitous in computer systems. Almost every language has a floating-point datatype; computers from PCs to supercomputers have floating-point accelerators; most compilers will be called upon to compile floating-point algorithms from time to time; and virtually every operating system must respond to floating-point exceptions such as overflow. This paper presents a tutorial on those aspects of floating-point that have a direct impact on designers of computer systems. It begins with background on floating-point representation and rounding error, continues with a discussion of the IEEE floating-point standard, and concludes with numerous examples of how computer builders can better support floating-point.

Categories and Subject Descriptors: (Primary) C.0 [Computer Systems Organization]: General -- *instruction set design*; D.3.4 [Programming Languages]: Processors -- *compilers, optimization*; G.1.0 [Numerical Analysis]: General -- *computer arithmetic, error analysis, numerical algorithms* (Secondary)

D.2.1 [Software Engineering]: Requirements/Specifications -- *languages*; D.3.4 Programming Languages]: Formal Definitions and Theory -- *semantics*; D.4.1 Operating Systems]: Process Management -- *synchronization*.

General Terms: Algorithms, Design, Languages

Additional Key Words and Phrases: Denormalized number, exception, floating-point, floating-point standard, gradual underflow, guard digit, NaN, overflow, relative error, rounding error, rounding mode, ulp, underflow.

# Neat Tool: “Herbie”

Some researchers developed a tool to automatically rearrange expressions to reduce FP error: <http://herbie.uwplse.org/demo/>



<http://herbie.uwplse.org/demo/>

**Herbie web demo**

See [the main page](#) for more info on Herbie.

Enter a formula below, hit Enter, and Herbie will try to improve it.