

# Numerical Linear Algebra – Pivoting and FLOP counting

## CS370 Lecture 30 – March 27, 2017

# Gaussian Elimination: $Ax = b$

(Some) numerical algorithms use Gaussian elimination, too. But it is interpreted differently...

Our view will be the following:

1. **Factor** matrix  $A$  into  $A = LU$ , where  $L$  and  $U$  are *triangular*.
2. **Solve**  $Lz = b$  for intermediate vector  $z$ .
3. **Solve**  $Ux = z$  for  $x$ .

(Later: We may also want to reorder (*permute*) the equations, which leads to the factorization  $PA = LU$ .)

# Numerical Problems – Factorization

During factoring, what if a diagonal entry,  $a_{k,k}$  is zero? Or close to zero?

For  $k = 1, \dots, n$

For  $i = k + 1, \dots, n$

$mult := a_{ik} / a_{kk}$

$a_{ik} := mult$

For  $j = k + 1, \dots, n$

$a_{ij} := a_{ij} - mult * a_{kj}$

EndFor

EndFor

EndFor

Divide by (exactly or nearly) zero?

# Numerical Problems

A division by zero is obviously bad news.

What about nearly zero?

This will make the multiplicative factor,  $a_{i,k}/a_{k,k}$ , **large** in magnitude.

A large factor can cause large floating point error during subtraction and magnify existing floating point error.

Leads to numerical instability!



Okay, who divided by zero?

# Solution: Row/Partial Pivoting

In earlier classes, you likely swapped rows if a diagonal entry was zero. This is called “row pivoting” or “partial pivoting”.

For numerical algorithms, one difference: we will **always** swap, to minimize floating point error incurred.

Strategy: Find the row with the *largest magnitude entry* in the current column beneath the current row, and swap those rows if larger than the current entry.

# Row/Partial Pivoting

Strategy: Find the row with the *largest magnitude entry* in the current column beneath the current row, and swap those rows if larger than the current entry.

e.g. 
$$\begin{bmatrix} 0 & 2 & 3 \\ 1 & 4 & 2 \\ -3 & 2 & -1 \end{bmatrix}$$

We immediately swap the first and third rows, since  $|-3| > |1| > 0$ .

# Pivoting with a Permutation Matrix

How can our factorization view account for row swaps?

Find a modified factorization of  $A$  such that  $PA = LU$  where  $P$  is a ***permutation matrix***.

A permutation matrix  $P$  is a matrix whose effect is to swap rows of the matrix it is applied to (i.e., multiplied with).

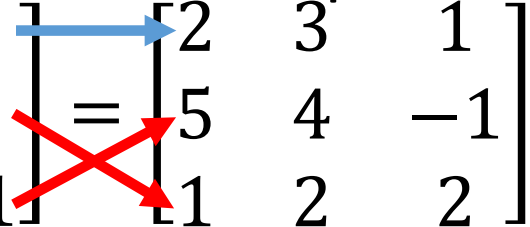
# Permutation Matrix

$P$  is simply a permuted (row-swapped) version of identity matrix,  $I$ .

e.g. to swap rows 2 and 3 of a 3x3 matrix, swap rows 2 and 3 of  $I$ .

$$P = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

Multiplying with a matrix performs the same swap:

$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} 2 & 3 & 1 \\ 1 & 2 & 2 \\ 5 & 4 & -1 \end{bmatrix} = \begin{bmatrix} 2 & 3 & 1 \\ 5 & 4 & -1 \\ 1 & 2 & 2 \end{bmatrix}$$




# Solving with the new factorization

Given the factorization  $PA = LU$ , multiplying  $Ax = b$  by  $P$  shows that  $PAx = Pb$ .

Using our factorization to replace  $PA$ , we have  $LUx = Pb$ .

Solve by first permuting entries of  $b$  according to  $P$ .

Then forward and backward substitution lets us find  $x$  as usual.

How do we determine  $P$  during factorization?

# Finding the permutation matrix, $P$

Start with  $P$  set to be an  $n \times n$  identity matrix,  $I$ .

Whenever we swap a pair of rows during LU factorization, also swap the corresponding rows of  $P$  (*including* the already stored factors).

The final  $P$  will be the desired permutation matrix.

# Example with pivoting

Let's try an example!

$$\begin{bmatrix} 1 & 4 & 5 \\ -2 & 3 & 3 \\ 3 & 0 & 6 \end{bmatrix} x = \begin{bmatrix} 4 \\ 1 \\ -3 \end{bmatrix}$$

# Summary:

## Gaussian Elimination with Row Pivoting

1. Perform LU factorization on  $A$  to find  $P$ ,  $L$ , and  $U$ .
2. Solve  $Lz = Pb$  for  $z$ . (Forward Solve.)
3. Solve  $Ux = z$  for  $x$ . (Backward Solve.)

Matlab's built-in LU factorization is the command:  
 $[L,U,P] = \text{lu}(A);$

# Costs of Gaussian Elimination

We want to know the (asymptotic) cost to solve a system of size  $n \times n$ .

We will measure cost in total *FLOPs: Floating point Operations*.

Approximate as the number of: *adds+subtracts+multiplies+divides*.

A true operation count would be hardware-dependent.


- e.g. Fused-multiply-add (FMA) may be a single operation.

(Careful: FLOPS is also floating-point-operations-*per-second*.)

# Cost of Factorization

```
For  $k = 1, \dots, n$   
  For  $i = k + 1, \dots, n$   
     $mult := a_{ik} / a_{kk}$   
     $a_{ik} := mult$   
    For  $j = k + 1, \dots, n$   
       $a_{ij} := a_{ij} - mult * a_{kj}$   
    EndFor  
  EndFor  
EndFor
```

2 FLOPs (1 subtraction, 1 multiply)  
in the innermost loop.



Summing over all the loops we get:

$$\sum_{k=1}^n \sum_{i=k+1}^n \sum_{j=k+1}^n 2 = \frac{2n^3}{3} + O(n^2)$$

The above requires using the following  
sum identities...

$$\sum_{i=1}^n i = \frac{n(n+1)}{2}$$

and

$$\sum_{i=1}^n i^2 = \frac{n(n+1)(2n+1)}{6}$$